Varied surfaces: further fitting removes most of the BC–DAgger fidelity gap without improving progress

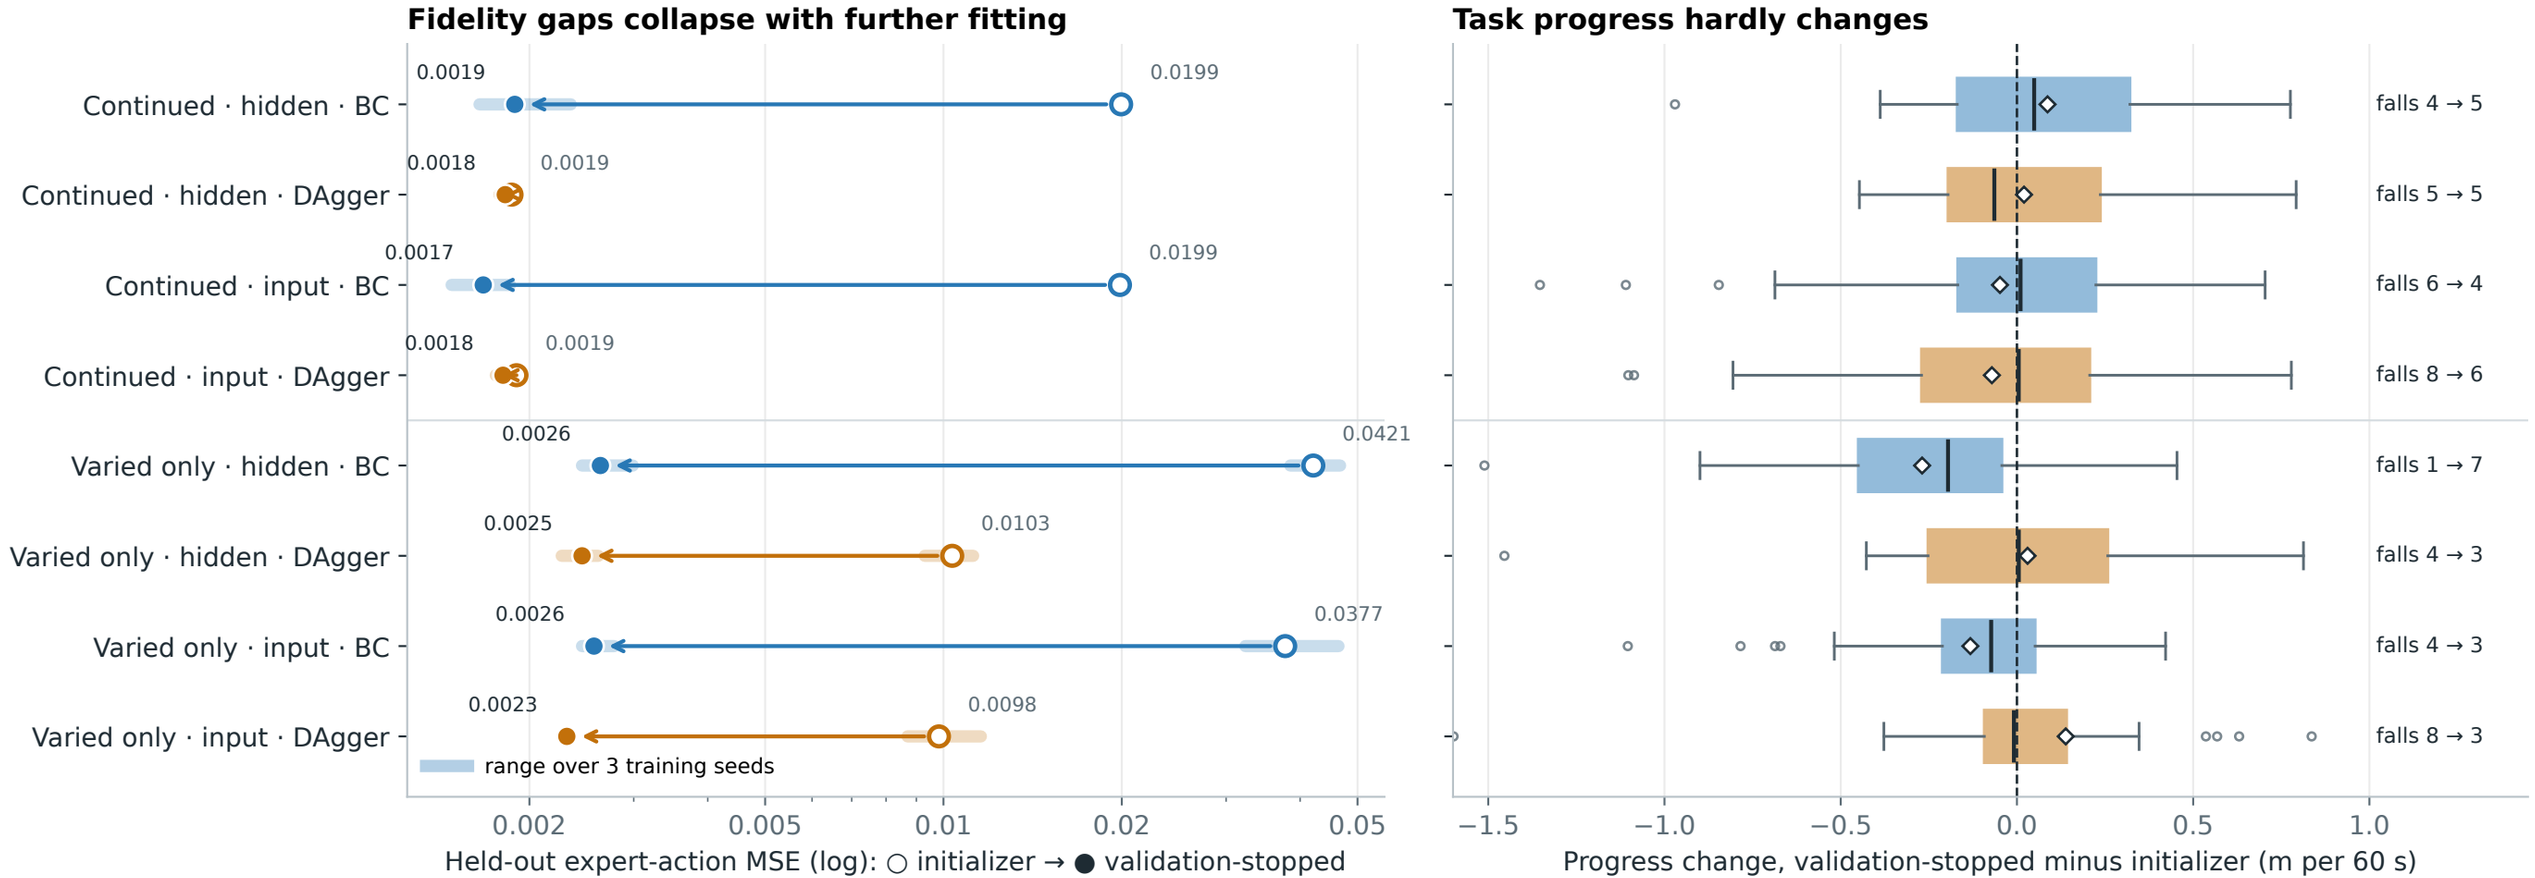

Each row: 3 training seeds. Initializers: continued = previous varied models; varied-only BC = 200-epoch control; varied-only DAgger = round-7 scratch model. Right: one value per sampled round (30), the seed-averaged paired change; box = IQR, ◇ = mean. Falls summed over the 3 seeds × 30 runs (90 active minutes). Re-running identical weights gave per-round SDs of 0.78–0.88 m, so family-mean changes smaller than ~0.2 m are within run-to-run noise.