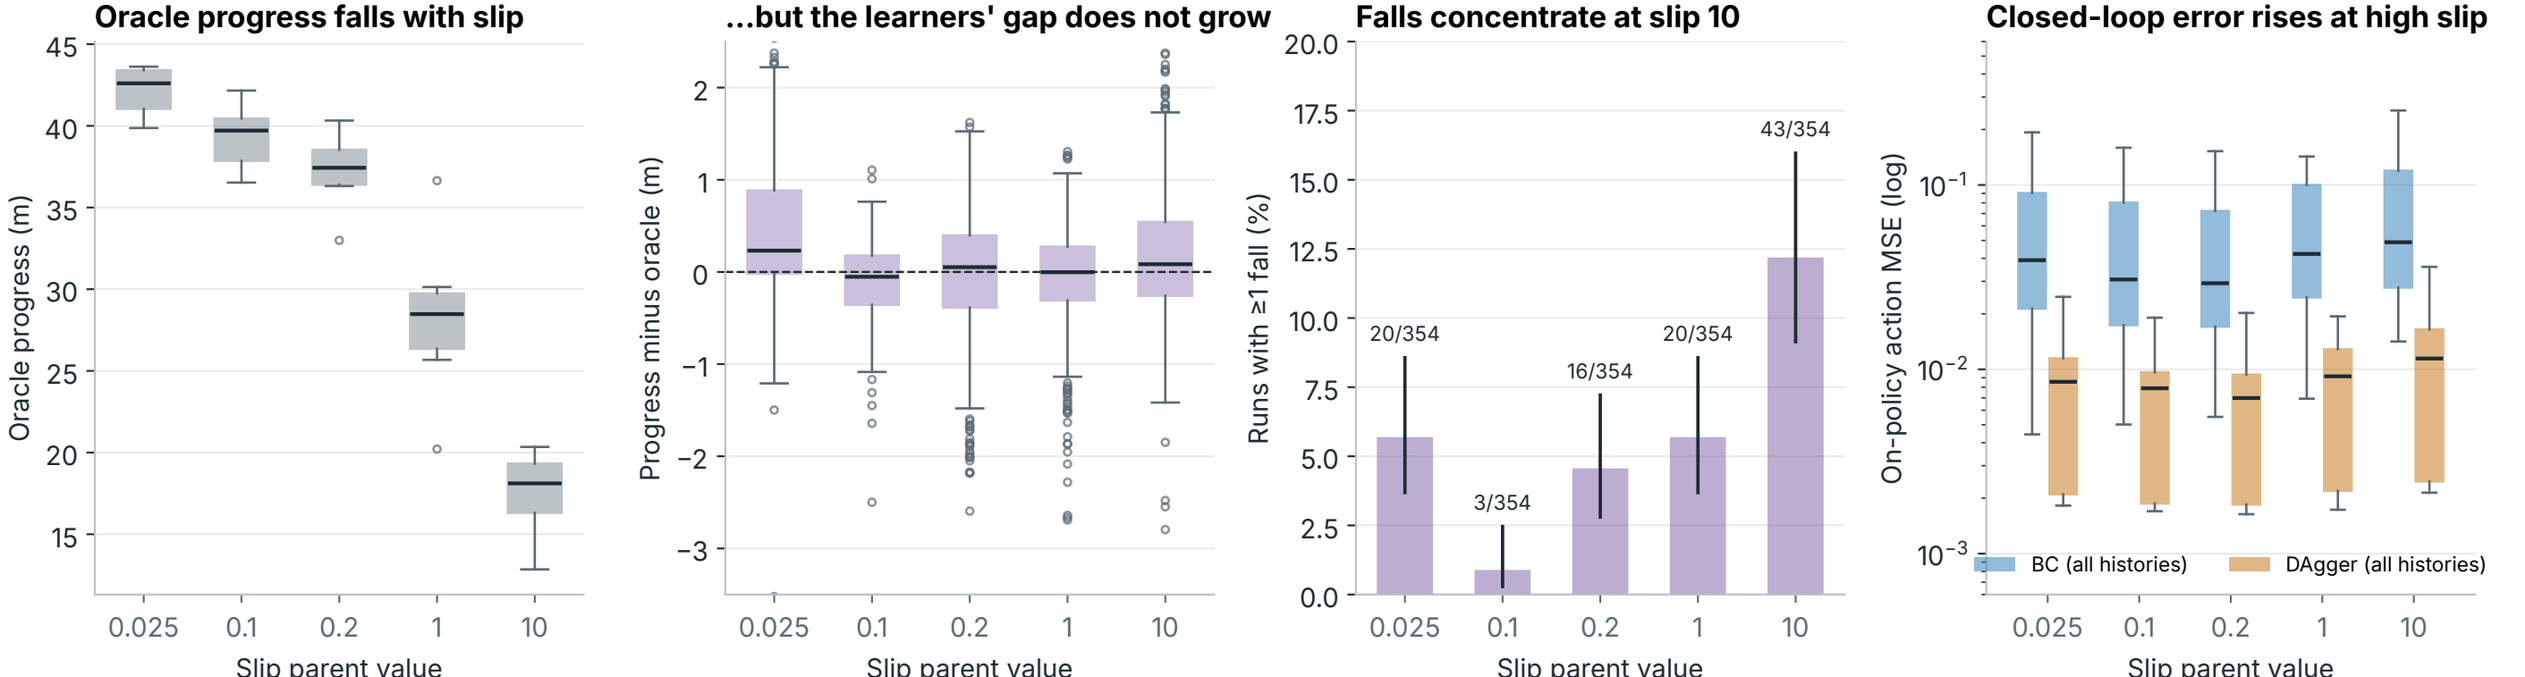

Each slip level holds 6 sampled rounds. Panels 2–4 pool every viable trained controller (BC, DAgger, GAIL-BC, GAIL-DM; all histories, both context types; 354 runs per level) — per-cohort values are in tables/gap_by_slip_level.csv. Fall interval: Wilson 95%, treating runs as independent.